
Structural Motifs Documentation


Release 0.3

Peter DeFord

Feb 04, 2019

CONTENTS

1	Contents	3
1.1	Introduction to Structural Motifs	3
1.2	Installation Instructions	4
1.3	Examples	5
1.4	Documentation	12
2	Indices and tables	13



StruM

Structural Motifs

Structural Motifs provide a way to represent DNA motifs based on the shape of the DNA instead of the DNA sequence. These motifs may represent the preferences driving specificity of transcription factors for their binding sites. For more information about the underlying model of StruMs, refer to [Introduction to Structural Motifs](#).

The package described here provides a framework for computing and working with StruMs, based on DNA shape features from the [Dinucleotide Property DataBase](#). Currently available capabilities include:

- Train models from known binding sites
- Use expectation maximization to find *de novo* motifs
- Provide utility to incorporate additional quantitative features, e.g. DNase hypersensitivity.
- Score matches of the motif to novel sequences

CONTENTS

1.1 Introduction to Structural Motifs

It has long been observed that transcription factors show a higher affinity to some sequences than others, and seem to tolerate some variability around these highest-affinity sequences.

The interaction between sequence specific transcription factors and DNA requires compatability at a complex interface of both electrostatics and 3D structure. Analagous to sequence representations of transcription factor binding sites, we assume that transcription factors have a preference for a specific shape of the DNA across the binding site and will tolerate some variation around that preferred shape.

Representing the values that a given position-specific shape feature may adopt as v_j , we expect the preference of the transcription factor to be normally distributed:

$$v_j \sim \mathcal{N}(\mu_j, \sigma_j^2)$$

In which case, given a set D of n sequences with m position-specific features that describe a sample of the binding sites preferred by a given transcription factor:

$$D = \left\{ \begin{pmatrix} v_{11} & v_{12} & \dots & v_{1m} \\ v_{21} & v_{22} & \dots & v_{2m} \\ \dots & \dots & \dots & \dots \\ v_{n1} & v_{n2} & \dots & v_{nm} \end{pmatrix} \right\}$$

we can compute a set of parameters ϕ describing the specificity of that transcription factor:

$$\phi = \left\{ \begin{pmatrix} \mu_1, \sigma_1 \\ \mu_2, \sigma_2 \\ \dots \\ \mu_m, \sigma_m \end{pmatrix} \right\}$$

If we also assume that each feature and each position is independent, then calculating the score s for the i th sequence becomes:

$$s_i = \prod_{j=1}^m P(v_{ij} | \mu_j, \sigma_j^2)$$

In order to avoid underflow issues during computation, all computations are done in log space.

1.1.1 *De novo* Motif finding

For *de novo* motif finding, an expectation maximization approach is employed. This approach assumes that there is exactly one occurrence of the binding site on each of the training sequences. This based on the OOPS model employed by MEME.

E-step

The likelihood (l_{ij}) of the j th position in the i th sequence being the start of the binding site is taken to be the score of the StruM at that position multiplied by the likelihood of the flanking regions matching the background model (ϕ_B):

$$l_{ij} = \prod_{n=1}^{j-1} P(v_{in}|\phi_B) \prod_{n=j}^{j+k-1} P(v_{in}|\phi_{i-j+1}) \prod_{n=j+k}^N P(v_{in}|\phi_B)$$

The likelihoods are then normalized on a by-sequence basis to produce M , the matrix of expected start positions:

$$M_{ij} = \frac{l_{ij}}{\sum_{j'=1}^m l_{ij'}}$$

M-step

The maximization step takes these likelihoods and calculates maximum likelihood values for μ and σ for each of the m position-specific features:

$$\mu_j = \frac{\sum_{i=1}^n \sum_v v_{ij} \cdot M_{ij}}{\sum_i \sum_j M_{ij}}$$
$$\sigma_j = \frac{\sum_{i=1}^n \sum_v \frac{(v_{ij} - \mu_j)^2 \cdot M_{ij}}{\sum_i \sum_j M_{ij} - \frac{\sum_i \sum_j M_{ij}^2}{\sum_i \sum_j M_{ij}}}}$$

1.2 Installation Instructions

1.2.1 Manual Installation of Dependencies

The StruM package depends on the following libraries:

- [numpy](#)
- [scipy](#)
- [matplotlib](#)
- [cython](#)
- [openmp](#)

1.2.2 Installing dependencies with Conda

By far the easiest way to install all of the dependencies is with [conda](#). With conda installed, the dependencies can be installed with:


```
conda install numpy scipy matplotlib cython openmp
```

1.2.3 Installing dependencies with PyPI

Another method for installing the dependencies manually is using the pip command from [PyPI](#). Then the dependencies can be installed with:

```
pip install numpy scipy matplotlib cython openmp
```

1.2.4 Installation of StruM Package

To install the StruM package, simply download the source code, navigate to that directory, and run the installation:

```
wget "https://github.com/pdeford/StructuralMotifs/archive/master.zip"
unzip -q master.zip
cd StructuralMotifs-master
```

If you only want it installed into your current directory for small scale use, use the following command:

```
python setup.py build_ext --inplace
```

If you want to install it system wide, use this command instead:

```
python setup.py build_ext
python setup.py install
```

1.3 Examples

This page contains simple examples showcasing the basic functionality of the StruM package.

Contents

- *Examples*
 - *Basic Example*
 - *De Novo Motif Finding Example*
 - *Adding Additional Features*
 - *Plotting a graphical representation of the StruM*

1.3.1 Basic Example

The basic usage of the StruM package covers the case where the binding site is already known. If the PWM is known already, a tool such as [FIMO](#) could be used to determine possible matches to the binding site.

In this case, a maximum likelihood StruM can be computed directly from these sequences. Source code for this example can be found here: [basic.py](#).

```
#!/usr/bin/env python

# Imports
import numpy as np
from strum import strum

# Sequences representing some of the variability of the
# FOXA1 binding site.
training_sequences = [
    "CTGTGCAAACA",
    "CTAAGCAAACA",
    "CTGTGCAAACA",
    "CTGTGCAAACA",
    "CAGAGCAAACA",
    "CTAAGCAAACA",
    "CTGTGCAAACA",
    "CAATGTAAACA",
    "CTGAGTAAATA",
]

# Initialize a new StruM object, using the basic features
# from the DiProDB table.
motif = strum.StruM(mode='basic')

# Use the training sequences to define the StruM,
# ensuring that the variation of all position-specific
# features is at least 10e-5 (lim)
motif.train(training_sequences, lim=10**-5)

# Print out what the PWM would look like for these sequences
motif.print_PWM(True)

# Define an example sequence to analyze.
test_sequence = "ACGTACTGCAGAGCAAACAACTGATCGGATC"
# Reverse complement it, as the best match may not be on
# the forward strand.
reversed_test = motif.rev_comp(test_sequence)

# Get a score of the similarity for each kmer in the test
# sequence to the StruM.
forward_scores = motif.score_seq(test_sequence)
reverse_scores = motif.score_seq(reversed_test)

# Find the best match.
## Determine the strand the best match lies on
if np.max(reverse_scores) > np.max(forward_scores):
    score, seq = reverse_scores, reversed_test
    strand = "-"
else:
    score, seq = forward_scores, test_sequence
    strand = "+"
## Find position in the sequence
best_pos = np.argmax(score)
best_seq = seq[best_pos : best_pos + motif.k]
best_score = score[best_pos]

# Print the results to screen
print "\n", strand, best_pos, best_seq, best_score
```

You should see the position weight matrix that would be derived from these sequences followed by the strand, position, matching kmer, and score for the best match in the test sequence to the StruM. The labels surrounding the PWM can be toggled off by changing the `labels` parameter in `strum.StruM.print_PWM()`.

```

      1      2      3      4      5      6      7      8      9     10     11
A 0.000 0.222 0.333 0.444 0.000 0.000 1.000 1.000 1.000 0.000 1.000
C 1.000 0.000 0.000 0.000 0.000 0.778 0.000 0.000 0.000 0.889 0.000
G 0.000 0.000 0.667 0.000 1.000 0.000 0.000 0.000 0.000 0.000 0.000
T 0.000 0.778 0.000 0.556 0.000 0.222 0.000 0.000 0.000 0.111 0.000

+ 8 CAGAGCAAACA -43.2745526744

```

1.3.2 De Novo Motif Finding Example

More often, the task is to take a set of sequences and perform *de novo* motif finding, simultaneously modeling the motif and identifying the putative binding sites. The StruM package allows you to do this using expectation maximization.

The following example uses the EM capabilities of the StruM package to identify the motif GATTACA randomly inserted into 50 sequences. The source code for this example can be downloaded here: [em.py](#).

Refer to `strum.StruM.train_EM()` for more info on the parameters for the EM algorithm.

```

#!/usr/bin/env python

# Imports
import numpy as np
from strum import strum

# 50 example sequences that contain the sequence GATTACA
training_sequences = [
    "CGATTACAGATCTCCGCGACCCCTT", "GATGATTACAAGATGCGTCGAATAT",
    "AGCCCTGTCCGCGAGATTACAACCAC", "AGTTAACTCCCTAGATTACATTTGT",
    "CGCTACAAAGTAAAGGAGATTACAT", "CTGATTACATCCTGTCCGAGGCGTG",
    "CGATTACATTCAACTAGATGCGCGC", "GAACGGCATGGGCGATTACAACACT",
    "CGGGGTGATCGTAATGTGATTACAA", "TGGTGCCCTGATTACACCTTACATG",
    "GGCGAAGATTACATCCTCGGCCCAT", "ATCGGATTACATGTACTCGTCCACG",
    "TCGGGGATTACAGGGGAGACGCTTA", "AGGTAGATTACATCGTTTTATTAGT",
    "TATTGTGCTCGATTACAAGCAGGCC", "CAGACCGCTTACACGTTGATTACAA",
    "GACACCCTCGATTACACCTCGTATA", "GGAACCGCGCGGATTACACGCGAGA",
    "GTTGATTACAAGGGAAACATACTTG", "CCCACACATTAGCTCGAGATTACAT",
    "GCAGAGTACCCTGCGGCGATTACAA", "ATACTCACGCATACAGATTACAAGA",
    "GGTATGCATCGCGATTACAGCACTG", "GGATTACAGTGAGCCTGCACCTTGA",
    "TTGGATTACATGGCCAACTCCACT", "GGCCGGCAGAGATTACACTAGAGAG",
    "ACAGATTACAGTGCAAATTGAGCAG", "CCACTGCATGACTGGATTACAGGCA",
    "CATGCCGGCGGTTAACGGATTACAC", "CCGATTACAAGTGCTCTGCACGGCG",
    "CATATAGAGGCGATTACAGCGTATC", "GGAACGATTACAGTGAGACTGCTCC",
    "ATGATTACAGCGAAACGTATTCAAA", "TTTTCGGATGATTACACATTCTTCT",
    "GTACAATGCATCGCGATTACAACAC", "GGATTACAAGTATCTGCCTGGATAC",
    "CTCCCGATTACATCAGGTACGTCCT", "TAGAGAAGATTACAGCCTACTATTG",
    "AAGCTTTGGCCGTACGATTACATC", "GTAAGATTACAAGTTCAGGGTGATC",
    "CATGATTACATTGGCGCCGACCTAC", "GCTGGATTACAATCATACCGTGTA",
    "GGTTAGGGATTACAAACAAGACGTG", "GACCGAGGTCTGATTACACTCCATC",
    "ATAGACGCGATTACAAGCACTCTAA", "TTTCCGTTCTGCAGCTGATTACAAC",
    "GGATTACACGCCTTCTCAAGCAGTG", "ATCCTAACAGGATTACAAGAATTAC",
    "TATGAAGCTGAAGAAGATTACAGCA", "CCTGTCTCAGATTACAGCACGGCGG",
]

```

```

# Initialize a new StruM object, using the DNA groove related
# features from the DiProDB table. Specify to use 4 cpus
# when doing EM.
motif = strum.StruM(mode='groove', n_process=4)

# Train the model on the training sequences using expectation
# maximization, ensuring that the variation of all
# position-specific features is at least 10e-5 (lim). Use
# a random_seed for reproducibility, and repeat with 20
# random restarts.
k = 8
motif.train_EM(training_sequences, fasta=False, k=k,
               lim=10**-5, random_seed=620, n_init=20)

# Examine the output, by identifying the best matching
# kmer in each of the training sequences.
out = []

for sequence in training_sequences:
    rseq = motif.rev_comp(sequence)
    s1 = motif.score_seq(sequence)
    s2 = motif.score_seq(rseq)
    i1 = np.argmax(s1)
    i2 = np.argmax(s2)
    if s1[i1] > s2[i2]:
        seq = sequence
        i = i1
        s = "+"
    else:
        seq = rseq
        i = i2
        s = "-"
    out.append(seq[i:i+k])
    print "{}{: <2} {: >{}} {} {}".format(
        s, i, seq[:i].lower(), len(seq)-k,
        seq[i:i+k].upper(), seq[i+k:].lower()
    )

# Summarize these best matches with a simple PWM.
nucs = dict(zip("ACGT", range(4)))
PWM = np.zeros([4,k])
for thing in out:
    for i,n in enumerate(thing):
        PWM[nucs[n], i] += 1

PWM /= np.sum(PWM, axis=0)
for row in PWM:
    print row

```

There are several sections that are reported in the output. First, the StruM gives a live update on how many iterations it took to converge for each of the random restarts. Once they have all completed, it shows the likelihoods for each of the restarts. The example output below highlights why it is important to repeat this process a number of times, as the results are highly variable. Finally we output a summary of the results. It is obvious that the model correctly identified the sequence NGATTACA as observed both in the highlighted sequences and the summary PWM at the bottom.

```

Retaining 50 out of 50 sequences, based on length (>25bp)
Detected cyclical likelihoods. Proceeding to max.
Converged after 10 iterations on likelihood

```

```

Converged after 5 iterations on likelihood
...
Restart Likelihoods: [1981.3168691613439, 1969.7041022864935, 1969.7041022864935, 650.
→89659719217332, 386.24238189057303, 369.62562708240256, 350.54071957020324, 350.
→54071957020324, 350.54071957020324, 327.8331936774872, 315.95426601483587, 242.
→73205512692093, 235.90375348128853, 149.92555284207441, -15.089679773020009, -68.
→060441249738645, -164.9202144304158, -373.49851273275601, -524.90598178843163, -524.
→91730067288267]
+0          CGATTACA gatctcccgcgaccctt
+2          ga TGATTACA agatgcgtcgaatat
+12      agccctgtccgc AGATTACA accac
+12      agttaactccct AGATTACA tttgt
...
[ 0.26  0.    1.    0.    0.    1.    0.    1. ]
[ 0.28  0.    0.    0.    0.    0.    1.    0. ]
[ 0.24  1.    0.    0.    0.    0.    0.    0. ]
[ 0.22  0.    0.    1.    1.    0.    0.    0. ]

```

1.3.3 Adding Additional Features

With the StruM package, in addition to the structural features provided in the DiProDB table, you can incorporate additional arbitrary features, as long as they provide a quantitative value across the binding site.

In the example below we will define a function that looks up a DNase signal from a bigwig file to incorporate. The file we will be using comes from a DNase experiment in K562 cells, mapped to *hg19* from the ENCODE project (ENCFF111KJD) and can be downloaded from [here](#).

The source code for this example can be found here: [DNase.py](#).

```

# Imports
import numpy as np
import sys

import bx.bbi.bigwig_file
from strum import strum

# Specify the path to where you downloaded the bigwig
# file. E.g. "/Users/user/Downloads/ENCFF111KJD.bigWig"
DNase_bigwig_path = sys.argv[1]

# Define the function to be used by the StruM when
# converting from sequence-space to structural-space.
# NOTE: This function takes additional parameters.
def lookup_DNase(seq, data, chrom, start, end):
    """Lookup the signal in a bigWig file, convert NaNs to
    0s, ensure strandedness, and return the modified signal.

    Parameters:
        data : (str) - Path to the bigWig file.
        chrom : (str) - Chromosome name.
        start : (int) - Base pair position on the chromosome of
            the beginning of the region of interest.
        end : (int) - Base pair position on the chromosome of
            the end of the region of interest.

    Returns:
        trace : (1D array) - The signal extracted from the
    """

```

```

                                bigWig file for the region of
                                interest.

    """
    # Open file
    bwh = bx.bbi.bigwig_file.BigWigFile(open(data))
    # Lookup signal for regions
    trace = bwh.get_as_array(chrom, min(start,end), max(start, end)-1)
    # Clean up NaNs
    trace[np.isnan(trace)] = 0.0
    # Ensure strandedness
    if start > end:
        trace = trace[::-1]
    return np.reshape(trace, [1,-1])

# Some example sequences and their chromosome positions,
# from human build hg19.
training_data = [
    ['GAGATCCTGGGTTTTCGAATCCCAGC', ('chr6', 26533165, 26533141)],
    ['GAGATCCTGGGTTTTCGAATCCCAGC', ('chr19', 33667901, 33667925)],
    ['GAGGTCCCGGGTTTCGATCCCCAGC', ('chr6', 28626033, 28626009)],
    ['GAGGTCCCTGGGTTTCGATCCCCAGT', ('chr6', 28763754, 28763730)],
    ['GGGGGCGTGGGTTTCGAATCCCACC', ('chr16', 22308468, 22308492)],
    ['GGGGGCGTGGGTTTCGAATCCCACC', ('chr5', 180614647, 180614671)],
    ['AAGGTCTGGGTTTCGAGCCCCAGT', ('chr11', 59318028, 59318004)],
    ['GAGGTCCCGGGTTTCAAATCCCGGA', ('chr1', 167684001, 167684025)],
    ['GAGGTCCCGGGTTTCAAATCCCGGA', ('chr7', 128423440, 128423464)],
]

# Initialize a new StruM object.
motif = strum.StruM(mode='basic')

print motif.features

# Update the StruM to incorporate the function
# defined above, drawing on the bigWig as the
# data source.
motif.update(data=DNase_bigwig_path, func=lookup_DNase,
             features=['k562_DNase'])

print motif.features

# Train the model using the modified StruM and the example
# data above
motif.train(training_data, lim=10.**-5)

# Evaluate the similarity to the model of a new sequence.
seq = 'GAGGTCCCGGGTTTCAATCCCCGGC'
position = ['chr2', 157257305, 157257329]

rseq = motif.rev_comp(seq)
rposition = [position[0], position[2], position[1]]

s = [
    motif.score_seq((seq, position)),
    motif.score_seq((rseq, rposition))
]

strands = ['+', '-']

```

```
print "Best match found on the '{}' strand".format(strands[np.argmax(s)])
```

This CTCF site is correctly identified as being in the forward orientation, and there is an additional feature being considered.

```
['Twist', 'Rise', 'Bend']
['Twist', 'Rise', 'Bend', 'k562_DNase']
Best match found on the '+' strand
```

1.3.4 Plotting a graphical representation of the StruM

You may find it useful to look at a graphical representation of the motif learned using the methods above. This can be accomplished using the `strum.StruM.plot()` method. Each of the features included in the StruM will be displayed as its own linegraph. The line represents the average value of that feature across the site, and the shaded area represents ± 1 standard deviation.


```
# Imports
from strum import strum

# Sequences representing some of the variability of the
# FOXA1 binding site.
training_sequences = [
    "CTGTGCAAACA",
    "CTAAGCAAACA",
    "CTGTGCAAACA",
    "CTGTGCAAACA",
    "CAGAGCAAACA",
    "CTAAGCAAACA",
    "CTGTGCAAACA",
    "CAATGTAAACA",
    "CTGAGTAAATA",
]

# Train the motif using the 'groove' related features
motif = strum.StruM(mode='groove')
motif.train(training_sequences, fasta=False)

# Plot the the StruM, and save to the specified filename
motif.plot('FOXA1_strum.png')
```

This code produces the image below. If you consider the Major Groove Distance, you will notice that there is more variation at the beginning of the motif (to the left) than at the end, as indicated by the shaded region. It is also clear to see that FOXA1 prefers a higher value at position 6, and a low value at position 9.



1.4 Documentation

1.4.1 StruM: Structural Motifs

This package provides functionality for computing structural representations of DNA sequence motifs. Estimates for DNA structure comes from the DiNucleotide Property Database (<http://diprodb.leibniz-flf.de/>).

This version relies on the Cython framework for speed purposes. Improvements in speed are particularly seen with scoring longer sequences.

`strum.rev_comp(str seq) → str`

Reverse complement a DNA sequence.

Parameters `seq(str.)` – A DNA sequence.

class `strum.StruM`

StruM Object: Train and work with Structural Motifs

StruMs can be learned via maximum likelihood (`train()`) from a known set of aligned binding sites, or via expectation maximization (`train_EM()`) from a set of broader regions.

For speed, this package will trim the sequences to be the same length. This can be user specified, else the shortest sequence will be used as a guide.

Additional features other than those from DiProDB can be incorporated by using the `update()` method.

`__init__(self, str mode='full', int n_process=1, custom_filter=[])`

Create a FastStruM object.

Parameters

- **mode** (*str.*) – Defines which subset of available features in the DiProDB table to use. Choose from: ['basic', 'protein', 'groove', 'proteingroove', 'unique', 'full', 'nucs', 'custom']

MODE	Features
basic	Twist, Rise, Bend.
protein	(for DNA-protein complex) Roll, Twist, Tilt, Slide, Shift, Rise.
groove	Major Groove Width, Major Groove Depth, Major Groove Size, Major Groove Distance, Minor Groove Width, Minor Groove Depth, Minor Groove Size, Minor Groove Distance.
protein-groove	The union of the “protein” and “groove” filters.
unique	Filters the table for the first occurrence of each type of feature.
full	All available double stranded DNA features.
nucs	Adenine content, Guanine content, Cytosine content, Thymine content.
custom	Manually select desired features.

- **custom_filter** (*list of ints.*) – Specifies the indices of desired features from the DiProDB table.
- **n_process** (*int.*) – Number of threads to use. -1 uses all processors.

define_PWM (*self, list seqs, weights=None*)

Computes a position weight matrix from sequences used to train the StruM.

Parameters

- **seqs** (*list of str.*) – Training set, composed of gapless alignment of binding sites of equal length.
- **weights** (*1D array of floats.*) – Weights to associate with each of the sequences in *seqs* to use in learning the motif.

Returns None. Sets the position weight matrix *self.PWM* based on the weighted sequences.

filter (*self*)

Update StruM to ignore uninformative position-specific features.

Features with a high variance, i.e. non-specific features, do not contribute to the specificity of the StruM model. Filtering them out may increase the signal-to-noise ratio. The position-specific features are rank ordered by their variance, and a univariate spline is fit to the distribution. The point of inflection is used as the threshold for masking less specific features.

Once this method is run and the attribute *self.filter_mask* is generated, two additional methods will become available: *score_seq_filt()* and *eval_filt()*.

func_translate (*self, seq, **kwargs*)

Convert sequence from string to structural representation, with additional features.

Parameters

- **seq** (*(str, [args])*) – DNA sequence, all uppercase characters, composed of letters from set ACGTN.
- ****kwargs** – Additional keyword arguments required by *self.func*.

Returns Sequence in structural representation.

Return type 1D numpy array of floats.

plot (*self*, *save_path*)

Save a graphical representation of the StruM.

Generates an image displaying the trends of each of the features in the StruM. The line indicates the average (scaled) value for that feature at each position. Shading represents +/- 1 standard deviation.

Note: This generates one row per feature. If you are including many features, your plot may be excessively tall.

Parameters **save_path** (*str.*) – Filename to use when saving the image.

print_PWM (*self*, *labels=False*)

Pretty prints the PWM to `std_out`.

Parameters **labels** (*bool.*) – Flag indicating whether to print the PWM with labels indicating the position associated with each column, and the nucleotide associated with each row.

Returns Formatted position weight matrix suitable for display, or use in the MEME suite, e.g. Also prints the PWM to `std_out`.

Return type `str`.

read_FASTA (*self*, *fasta_file*)

Reads a FASTA formatted file for headers and sequences.

Parameters **fasta_file** (*file object.*) – FASTA formatted file containing DNA sequences.

Returns The headers and sequences from the FASTA file, as two separate lists.

Return type (`list`, `list`)

score_seq (*self*, *str seq*, ***kwargs*)

Scores a sequence using pre-calculated StruM.

Parameters **seq** (*str.*) – DNA sequence, all uppercase characters, composed of letters from set ACGTN.

Returns Vector of scores for similarity of each kmer in `seq` to the StruM.

Return type 1D array.

score_seq_filt (*self*, *seq*, ***kwargs*)

A variation on `score_seq()` that masks non-specific features.

Once the `self.filter_mask` is generated, this method becomes available. This scores a sequence with the precomputed StruM, masking non-specific features.

Refer to `score_seq()` for more information about the arguments.

text_strum (*self*, *avg*, *std*, *avg_range=None*, *std_range=None*, *colorbar=False*)

Gerates a ANSI colored Unicode text representation of a StruM.

Parameters

- **avg** (*Array of floats.*) – The average values of the StruM to plot.

- **std** (*Array of floats.*) – The standard deviation values of the StruM to plot.
- **avg_range** (*tuple (low_value, high_value).*) – Seven bins of even width will be generated across this range, and the avg values will be assigned to one of the bins.
- **std_range** (*tuple (low_value, high_value).*) – Eight bins of even width will be generated across this range, and the std values will be assigned to one of the bins.
- **colorbar** (*bool.*) – Whether to include a colorbar

Returns Returns a tuple of strings. The first element is the formatted StruM representation. If `colorbar=True`, the second element is the colorbar string.

Return type tuple of strings.

train (*self, training_sequences, weights=None, lim=None, **kwargs*)

Learn structural motif from a set of known binding sites.

Parameters

- **training_sequences** (*list of str. Or if updated, list of tuples: (str, [args])*) – Training set, composed of gapless alignment of binding sites of equal length.
- **weights** (*1D array of floats.*) – Weights to associate with each of the sequences in `training_sequences` to use in learning the motif.
- **lim** (*float*) – Minimum value allowed for variation in a given position-specific-feature. Useful to prevent *any* deviation at that position from resulting in a probability of 0.

Returns None. Defines the structural motif `self.strum` and the corresponding position weight matrix `self.PWM`, sets attribute `self.fit = True`.

train_EM (*self, data, fasta=True, params=None, int k=10, int max_iter=1000, double convergence_criterion=0.001, random_seed=0, int n_init=1, lim=None, seqlength=None, background=None, seed_motif=None, bool verbose=False*)

Performs Expectation-Maximization on a set of sequences to find motif.

Parameters

- **data** (*list of str, open file object referring to a FASTA file.*) – A set of sequences to use for training the model. Assumed to have one occurrence of the binding site per sequence.
- **fasta** (*bool.*) – Flag indicating whether data points to an open file object containing a FASTA formatted file with DNA sequences.
- **params** (**args, **kwargs.*) – Additional parameters to pass to `self.func`, if defined.
- **k** (*int.*) – Size of binding site to consider. Since dinucleotides are considered, in sequence-space the size of the binding site will be $k + 1$.
- **max_iter** (*int.*) – Maximum number of iterations of Expectation Maximization to perform if convergence is not attained.
- **convergence_criterion** (*float.*) – If the change in the likelihood between two iterations is less than this value, the model is considered to have converged.
- **random_seed** (*int.*) – Seed for the random number generator used in the EM algorithm for initialization.
- **n_init** (*int.*) – Number of random restarts of the EM algorithm to perform.

- **lim** (*float*) – Minimum value allowed for variation in a given position-specific-feature. Useful to prevent *any* deviation at that position from resulting in a probability of 0.
- **seqlength** (*int*.) – If set, the sequences in the training data will be trimmed symmetrically to this length. .. note:: This must be longer than the shortest sequence.
- **background** (*str*.) – Method to use for computing the background model. Default is to assume equal probability of each dinucleotide. Passing "compute" adapts the background model to represent the dinucleotide frequencies in the training sequences. Otherwise this can be the path to a tab delimited file specifying the representation of each dinucleotide. E.g.

```
AA  0.0625
AC  0.0625
...
```

- **seed_motif** (*StruM.strum*.) – Optional. A *StruM.strum* to use for initializing the Expectation-Maximization algorithm. If set, *k* will be replaced by the corresponding value from the *seed_motif*'s shape. *n_init* will also be reset to 1.
- **verbose** – Specifies whether to print a text version of the converging StruM at each iteration. Default *False*.

Returns *None*. Defines the structural motif *self.strum* and the corresponding position weight matrix *self.PWM*.

translate (*self, str seq, **kwargs*)

Convert sequence from string to structural representation.

Parameters

- **seq** (*str*.) – DNA sequence, all uppercase characters, composed of letters from set ACGTN.
- ****kwargs** – Ignored

Returns Sequence in structural representation.

Return type 1D numpy array of floats.

update (*self, features, func, data=None*)

Update the StruM to incorporate additional features.

Using this method will change the behavior of other methods, especially the *translate()* method is replaced by *func_translate()*.

Parameters

- **features** (*list of str*.) – Text description or label of the feature(s) being added into the model.
- **func** (*function*.) – The scoring function that produces the additional features. These may be computed on sequence alone, or by incorporating additional data. The output must be an array of shape *[n, l-1]*, where *n* is the number of additional features (*len(features)*) and *l* is the length of the sequence being scored. The first argument of the function must be the sequence being scored.
- **data** – Additional data that is used by the new function. May be a lookup table, for example, or a reference to an outside file.

INDICES AND TABLES

- `genindex`
- `modindex`
- `search`

PYTHON MODULE INDEX

S

strum, ??